

# Diabetes Project

January 29, 2024

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')
```

```
[2]: diabetes_data = pd.read_csv('health_care_diabetes.csv')
```

```
[3]: # Preview data
diabetes_data.head()
```

```
[3]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0             6     148             72             35         0  33.6
1             1      85             66             29         0  26.6
2             8     183             64              0         0  23.3
3             1      89             66             23        94  28.1
4             0     137             40             35       168  43.1
```

```
   DiabetesPedigreeFunction  Age  Outcome
0                0.627     50         1
1                0.351     31         0
2                0.672     32         1
3                0.167     21         0
4                2.288     33         1
```

```
[4]: # Dataset dimensions - (rows, columns)
diabetes_data.shape
```

```
[4]: (768, 9)
```

```
[5]: # Features data-type
diabetes_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64

dtypes: float64(2), int64(7)

memory usage: 54.1 KB

```
[6]: # Statistical summary
diabetes_data.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[8]: # Count of null values
diabetes_data.isnull().sum()
```

Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0

```
Age          0
Outcome      0
dtype: int64
```

```
[9]: diabetes_data.describe().T
```

```
[9]:
```

	count	mean	std	min	25%	\
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	
Glucose	768.0	120.894531	31.972618	0.000	99.00000	
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000	
Insulin	768.0	79.799479	115.244002	0.000	0.00000	
BMI	768.0	31.992578	7.884160	0.000	27.30000	
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	
Age	768.0	33.240885	11.760232	21.000	24.00000	
Outcome	768.0	0.348958	0.476951	0.000	0.00000	

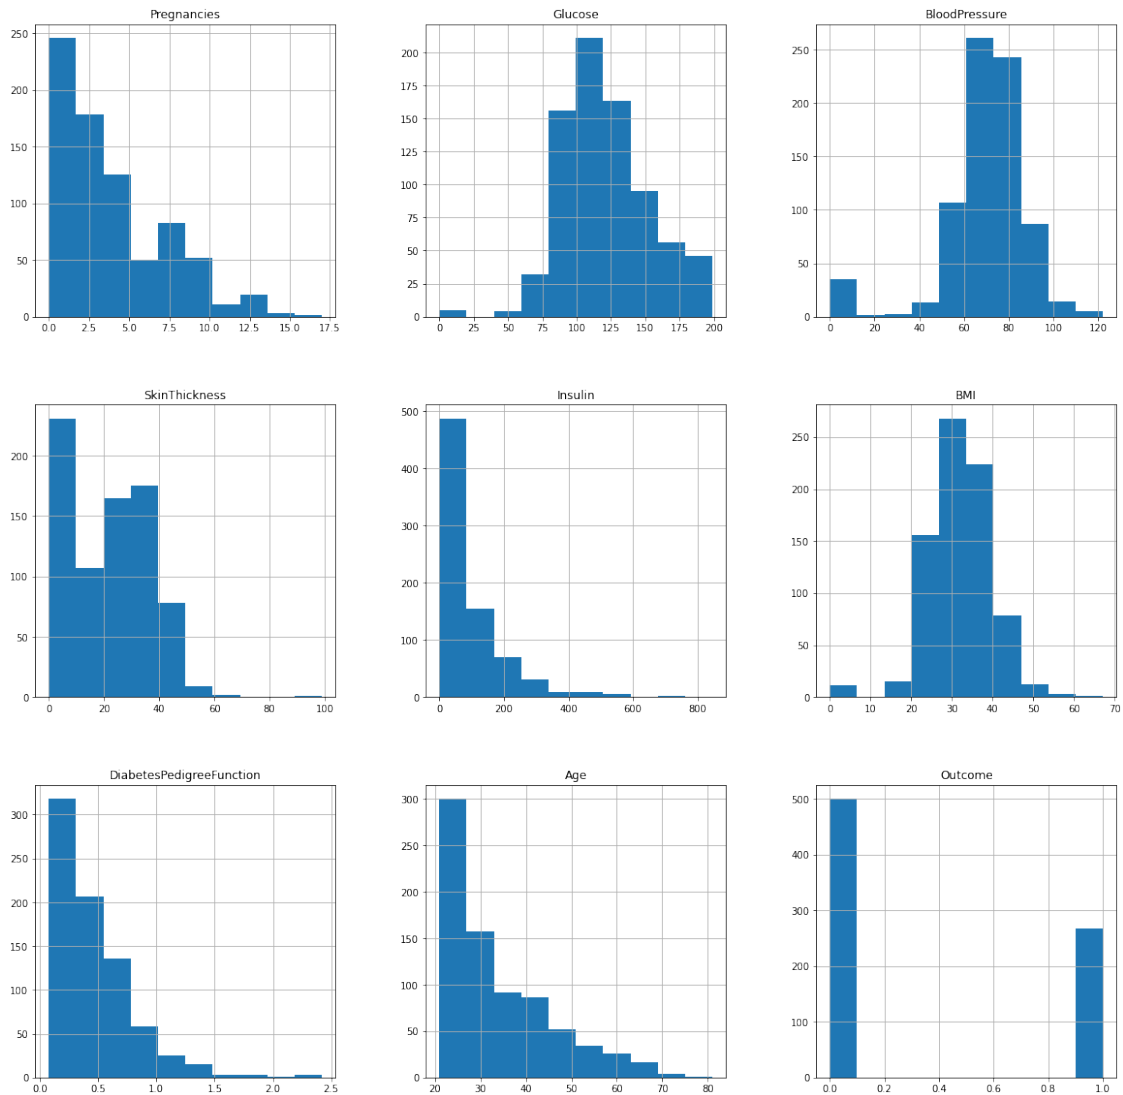
	50%	75%	max
Pregnancies	3.0000	6.00000	17.00
Glucose	117.0000	140.25000	199.00
BloodPressure	72.0000	80.00000	122.00
SkinThickness	23.0000	32.00000	99.00
Insulin	30.5000	127.25000	846.00
BMI	32.0000	36.60000	67.10
DiabetesPedigreeFunction	0.3725	0.62625	2.42
Age	29.0000	41.00000	81.00
Outcome	0.0000	1.00000	1.00

```
[10]: diabetes_data_copy = diabetes_data.copy(deep = True)
diabetes_data_copy[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']]
↳=
↳diabetes_data_copy[['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']]
↳replace(0,np.NaN)

## showing the count of Nans
print(diabetes_data_copy.isnull().sum())
```

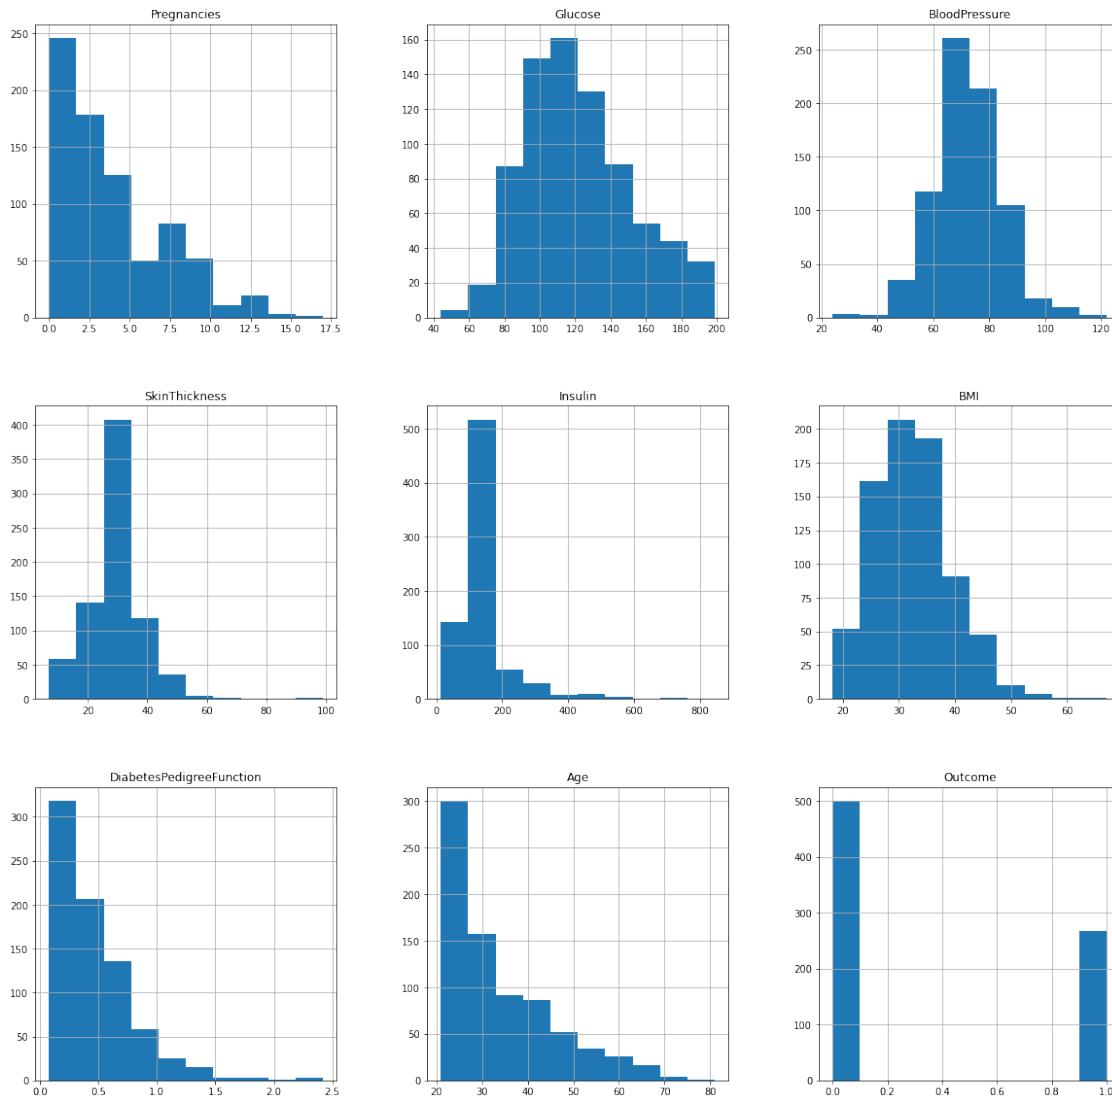
```
Pregnancies      0
Glucose           5
BloodPressure     35
SkinThickness     227
Insulin           374
BMI               11
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64
```

```
[11]: p = diabetes_data.hist(figsize = (20,20))
```



```
[12]: diabetes_data_copy['Glucose'].fillna(diabetes_data_copy['Glucose'].mean(),
      ↪inplace = True)
diabetes_data_copy['BloodPressure'].fillna(diabetes_data_copy['BloodPressure'].
      ↪mean(), inplace = True)
diabetes_data_copy['SkinThickness'].fillna(diabetes_data_copy['SkinThickness'].
      ↪median(), inplace = True)
diabetes_data_copy['Insulin'].fillna(diabetes_data_copy['Insulin'].median(),
      ↪inplace = True)
diabetes_data_copy['BMI'].fillna(diabetes_data_copy['BMI'].median(), inplace =
      ↪True)
```


```
[13]: p = diabetes_data_copy.hist(figsize = (20,20))
```




```
[14]: diabetes_data.shape
```

```
[14]: (768, 9)
```

```
[23]: sns.countplot(y=diabetes_data.dtypes ,data=diabetes_data)
plt.xlabel("count of each data type")
plt.ylabel("data types")
plt.show()
```




```
[24]: ## null count analysis
import missingno as msno
p=msno.bar(diabetes_data)
```

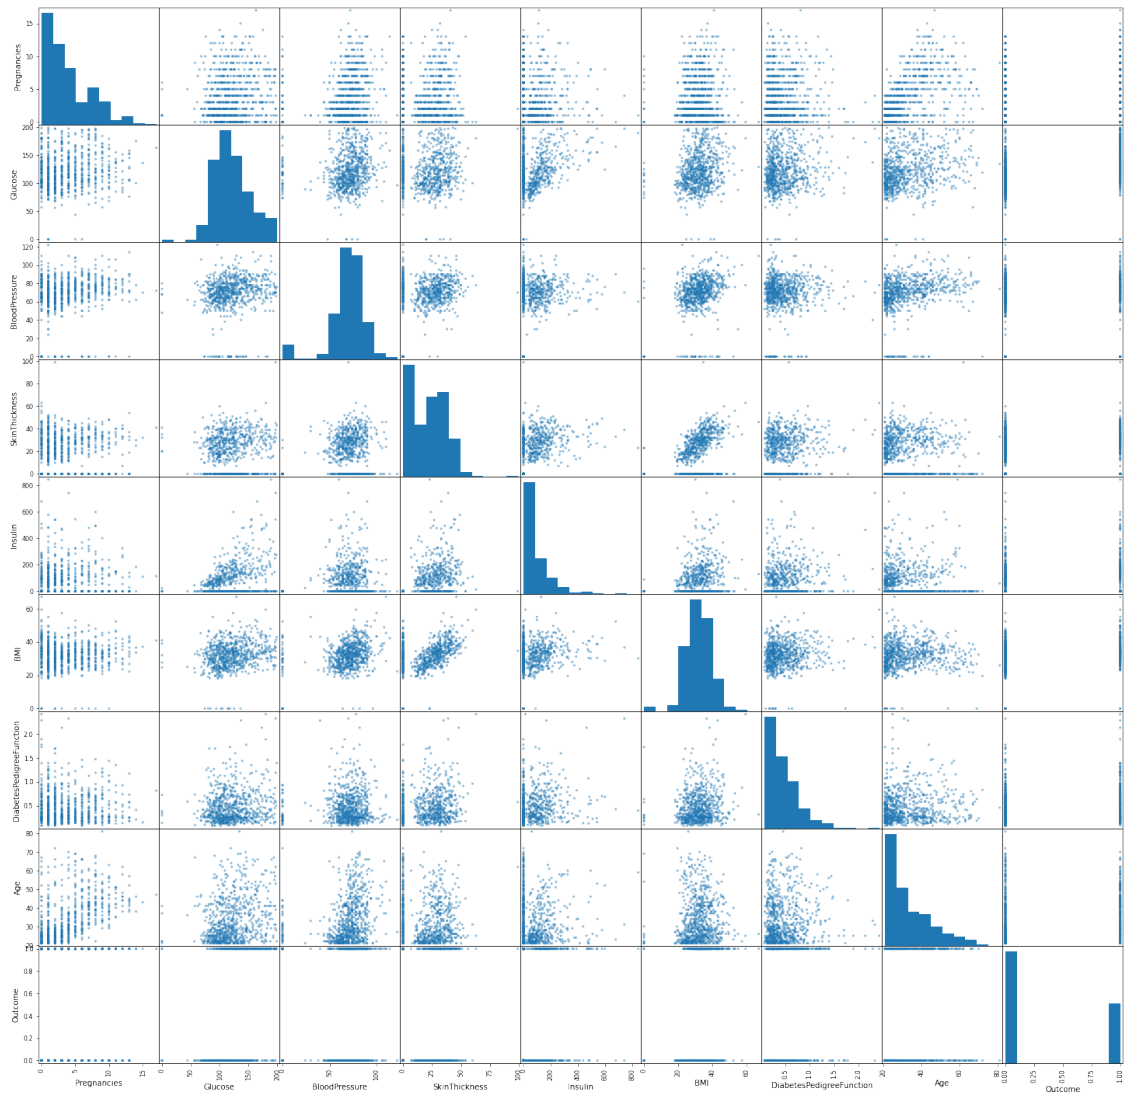


```
[25]: ## checking the balance of the data by plotting the count of outcomes by their
      ↪ value
      color_wheel = {1: "#0392cf",
                     2: "#7bc043"}
      colors = diabetes_data["Outcome"].map(lambda x: color_wheel.get(x + 1))
      print(diabetes_data.Outcome.value_counts())
      p=diabetes_data.Outcome.value_counts().plot(kind="bar")
```

```
0    500
1    268
Name: Outcome, dtype: int64
```

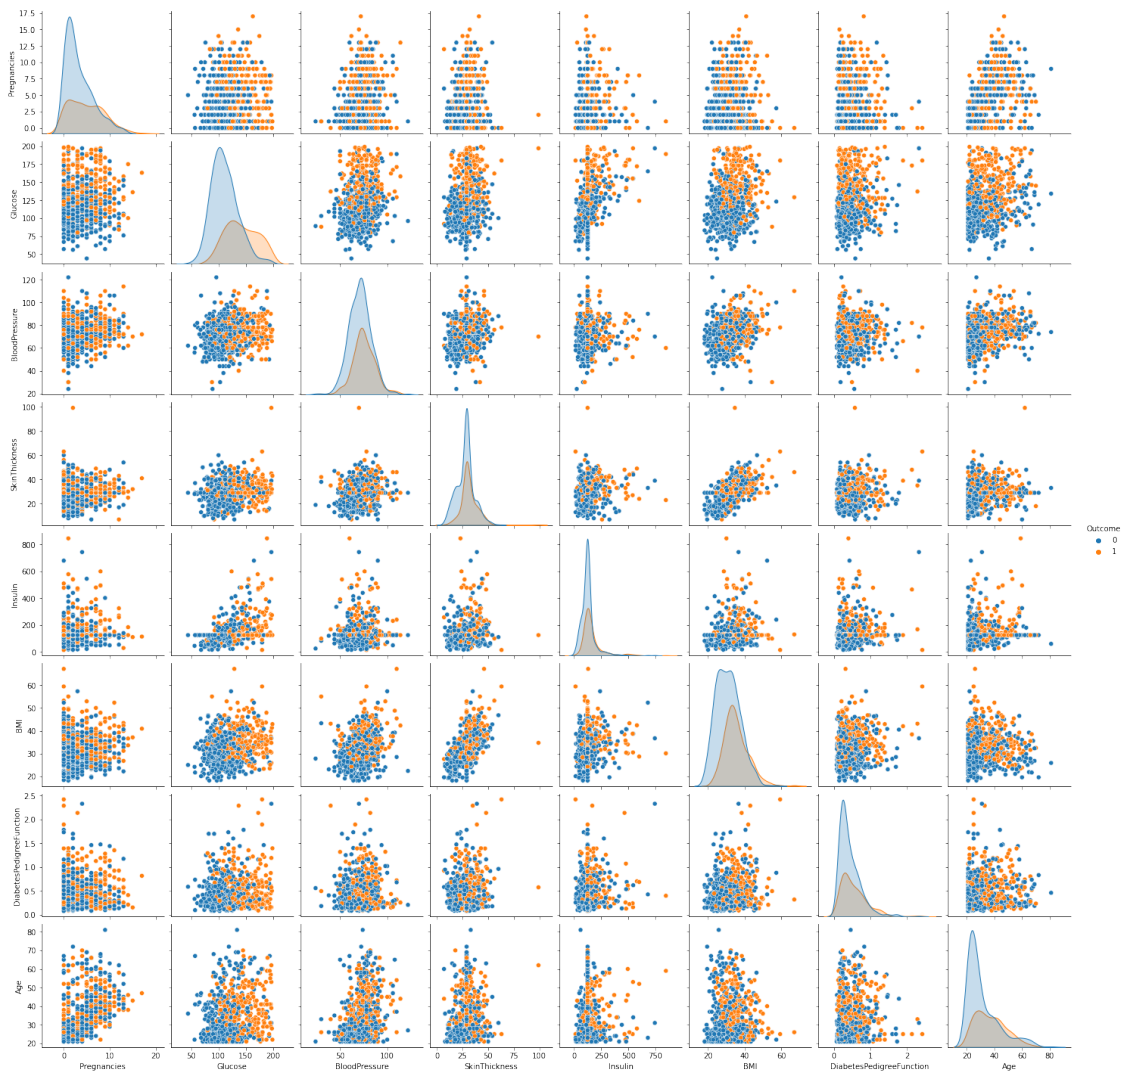


```
[27]: from pandas.plotting import scatter_matrix
      p=scatter_matrix(diabetes_data,figsize=(25, 25))
```




```
[28]: p=sns.pairplot(diabetes_data_copy, hue = 'Outcome')
```






```
[29]: plt.figure(figsize=(14,12))
p=sns.heatmap(diabetes_data.corr(), annot=True,cmap='RdYlGn')
```



```
[31]: plt.figure(figsize=(14,12))
p=sns.heatmap(diabetes_data_copy.corr(), annot=True,cmap = 'RdYlGn')
```



```
[32]: from sklearn.preprocessing import StandardScaler
sc_X = StandardScaler()
X = pd.DataFrame(sc_X.fit_transform(diabetes_data_copy.drop(["Outcome"],axis = 1)),
                 columns=['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
                 'Insulin',
                 'BMI', 'DiabetesPedigreeFunction', 'Age'])
```

```
[33]: X.head()
```

```
[33]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin    BMI \
0      0.639947  0.865108    -0.033518    0.670643 -0.181541  0.166619
1     -0.844885 -1.206162    -0.529859   -0.012301 -0.181541 -0.852200
2      1.233880  2.015813    -0.695306   -0.012301 -0.181541 -1.332500
3     -0.844885 -1.074652    -0.529859   -0.695245 -0.540642 -0.633881
```

```
4      -1.141852  0.503458      -2.680669      0.670643  0.316566  1.549303
```

```
DiabetesPedigreeFunction      Age
0      0.468492  1.425995
1     -0.365061 -0.190672
2      0.604397 -0.105584
3     -0.920763 -1.041549
4      5.484909 -0.020496
```

```
[34]: y = diabetes_data_copy.Outcome
```

```
[35]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=1/
↳3,random_state=42, stratify=y)
```

```
[36]: from sklearn.neighbors import KNeighborsClassifier
```

```
test_scores = []
train_scores = []

for i in range(1,15):

    knn = KNeighborsClassifier(i)
    knn.fit(X_train,y_train)

    train_scores.append(knn.score(X_train,y_train))
    test_scores.append(knn.score(X_test,y_test))
```


```
[37]: max_train_score = max(train_scores)
train_scores_ind = [i for i, v in enumerate(train_scores) if v ==
↳max_train_score]
print('Max train score {} % and k = {}'.
↳format(max_train_score*100,list(map(lambda x: x+1, train_scores_ind))))
```

Max train score 100.0 % and k = [1]

```
[38]: max_test_score = max(test_scores)
test_scores_ind = [i for i, v in enumerate(test_scores) if v == max_test_score]
print('Max test score {} % and k = {}'.
↳format(max_test_score*100,list(map(lambda x: x+1, test_scores_ind))))
```

Max test score 76.5625 % and k = [11]

```
[44]: plt.figure(figsize=(12,5))
p = sns.lineplot(train_scores,label='Train Score')
p = sns.lineplot(test_scores,label='Test Score')
```




```
[45]: knn = KNeighborsClassifier(11)
```

```
knn.fit(X_train,y_train)
knn.score(X_test,y_test)
```

```
[45]: 0.765625
```

```
[47]: import matplotlib
matplotlib.use('Agg')
from mlxtend.plotting import plot_decision_regions
import matplotlib.pyplot as plt
value = 20000
width = 20000
plot_decision_regions(X.values, y.values, clf=knn, legend=2,
                      filler_feature_values={2: value, 3: value, 4: value, 5:
↪value, 6: value, 7: value},
                      filler_feature_ranges={2: width, 3: width, 4: width, 5:
↪width, 6: width, 7: width},
                      X_highlight=X_test.values)

# Adding axes annotations
#plt.xlabel('sepal length [cm]')
#plt.ylabel('petal length [cm]')
plt.title('KNN with Diabetes Data')
plt.show()
```




```
[48]: #import confusion_matrix
from sklearn.metrics import confusion_matrix
#let us get the predictions using the classifier we had fit above
y_pred = knn.predict(X_test)
confusion_matrix(y_test,y_pred)
pd.crosstab(y_test, y_pred, rownames=['True'], colnames=['Predicted'],
↪margins=True)
```

```
[48]: Predicted    0    1  All
True
0          142   25  167
1           35   54   89
All         177   79  256
```

```
[49]: y_pred = knn.predict(X_test)
from sklearn import metrics
cnf_matrix = metrics.confusion_matrix(y_test, y_pred)
p = sns.heatmap(pd.DataFrame(cnf_matrix), annot=True, cmap="YlGnBu" ,fmt='g')
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
```

```
[49]: Text(0.5, 15.0, 'Predicted label')
```




```
[52]: #import classification_report
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.85	0.83	167
1	0.68	0.61	0.64	89
accuracy			0.77	256
macro avg	0.74	0.73	0.73	256
weighted avg	0.76	0.77	0.76	256

```
[53]: from sklearn.metrics import roc_curve
y_pred_proba = knn.predict_proba(X_test)[: ,1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
```

```
[54]: plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr,tpr, label='Knn')
plt.xlabel('fpr')
plt.ylabel('tpr')
```

```
plt.title('Knn(n_neighbors=11) ROC curve')
plt.show()
```



```
[55]: #Area under ROC curve
from sklearn.metrics import roc_auc_score
roc_auc_score(y_test,y_pred_proba)
```

```
[55]: 0.8193500639171096
```

```
[56]: #import GridSearchCV
from sklearn.model_selection import GridSearchCV
#In case of classifier like knn the parameter to be tuned is n_neighbors
param_grid = {'n_neighbors':np.arange(1,50)}
knn = KNeighborsClassifier()
knn_cv= GridSearchCV(knn,param_grid,cv=5)
knn_cv.fit(X,y)

print("Best Score:" + str(knn_cv.best_score_))
print("Best Parameters: " + str(knn_cv.best_params_))
```

```
Best Score:0.7721840251252015
Best Parameters: {'n_neighbors': 25}
```

```
[ ]:
```